

Signal Processing Tutorial

Research School: Mathematics, Signal Processing and Learning

Laurent Oudre

laurent.oudre@ens-paris-saclay.fr

CIRM, Marseille, France

January 27th, 2021

Contents

1. Introduction and data

- 1.1 Problem statement
- 1.2 Data

2. Basic tools and notions for signal processing

- 2.1 Discrete Fourier Transform
- 2.2 Stationarity, ergodicity and autocorrelation function
- 2.3 Spectrogram

3. Change-point detection

- 3.1 Cost functions
- 3.2 Search method

4. Convolutional dictionary learning

- 4.1 Problem
- 4.2 Alternated resolution

Contents

1. Introduction and data

1.1 Problem statement


1.2 Data

2. Basic tools and notions for signal processing

3. Change-point detection

4. Convolutional dictionary learning


Gait analysis



Why is it important to study locomotion?

- ▶ Most common dynamic human activity
- ▶ Can reveal a large number of neurological, orthopedic, rheumatological disorders...
- ▶ Strong influence on daily life : risk of falling, frailty, autonomy, dependency...

Gait analysis



How can we study locomotion?

- ▶ Early tests: clinical examination by the physician, functional tests, clinical questionnaires

+	Easy to perform, use of clinical expertise
-	Lack of precision, difficult to objectively compare two sessions

- ▶ Dedicated platforms for the study of locomotion: instrumented mats, video/optical systems

+	Great precision, extraction of a large number of useful features, objective quantification
-	Expensive, difficult to put in practice

Main principles

★ Objective quantification of human gait

→ Use of sensors and physiological measurements

★ Longitudinal follow-up and inter-individual comparison

→ Need for a fixed protocol

★ Experimentation outside the laboratory and on the field

→ User-mounted sensors and fully automatic device for consultation and routine use


★ Clean data

→ Control of the entire measurement chain, robust and reproducible algorithms

★ Willingness to capture the expertise of the clinician



→ Clinical annotations and metadata

Protocol



- ▶ Sequence of activities:
 - ▶ stand for 6 s,
 - ▶ walk 10 m at preferred walking speed on a level surface to a previously shown turn point,
 - ▶ turn around (without previous specification of a turning side),
 - ▶ walk back to the starting point,
 - ▶ stand for 2 s.
- ▶ Subjects walked at their comfortable speed with their shoes and without walking aid.

Sensors




- ▶ IMU (Inertial Measurement Unit) record linear accelerations (3D), angular velocities (3D) and magnetic fields (3D) on each foot
- ▶ Sensor frame consists of 3-axis (X , Y , Z)
- ▶ **For this tutorial: angular velocity around axis Y**

Database

221 recordings:

- ▶ Healthy subjects had no known medical impairment (labelled as "T" for Témoin).
- ▶ The orthopedic group is composed of 3 cohorts of distinct pathologies: lower limb osteoarthritis (ArtH, ArtG), cruciate ligament injury (LCA), knee injury (Genou)
- ▶ The neurological group is composed of 2 cohorts: cerebellar disorder (CER) and radiation induced leukoencephalopathy (LER)

Main scientific questions



Non-stationary signals

→ How can we detect the different regimes (stop, walking, U-turn...)?

Presence of repetitive patterns: the steps

→ What are they? How could we automatically extract them?

Robust feature extraction

→ How could we extract relevant features for longitudinal follow-up and inter-individual comparison?

Contents

1. Introduction and data

2. Basic tools and notions for signal processing

2.1 Discrete Fourier Transform

2.2 Stationarity, ergodicity and autocorrelation function

2.3 Spectrogram

3. Change-point detection

4. Convolutional dictionary learning

What is signal?

- ▶ A time series (or signal) is a series of data points indexed in time order
- ▶ In practice, array of real numbers of size $D \times N$ where D is the number of dimensions and N the number of samples
 - ▶ Sample number n

n	0	1	2	3	4	5	6
-----	---	---	---	---	---	---	---

- ▶ Time series values $x[n]$

$x[n]$	0.7	0.2	0.8	0.9	0.3	0.2	0.7
	0.4	0.1	0.6	0.2	0.5	0.6	0.3

- ▶ Time stamps $t[n]$

$t[n]$	16:30:01	16:30:23	16:31:43	16:32:38	16:33:06	16:33:16	16:33:56
--------	----------	----------	----------	----------	----------	----------	----------

Two visions: physics vs. statistics

- ▶ The notion of time have been used and modeled in physics since 18th century and before (eg. Fourier transform).

First vision : a time series $x[1 : N]$ is the result of the digitization of a physical phenomenon $x(t)$. Physical properties of this phenomenon can be retrieved and analyzed through the study of $x[1 : N]$ (and vice-versa).

- ▶ Randomness can also play a part to model a wider class of signals.

Second vision : a time series $x[1 : N]$ is a realization of a stochastic process $X[1 : N]$. Statistical properties of this phenomenon can be retrieved and analyzed through the study of $x[1 : N]$ (and vice-versa).

In most cases, both approaches can be combined.

Some useful signal processing tools

In the following, we will introduce basic signal processing tools and apply them to our signals:

- ▶ Discrete Fourier Transform (DFT)
- ▶ Notion of stationarity, ergodicity and autocorrelation function
- ▶ Spectrogram

Sampling and Fourier analysis

- ▶ Most tools for signal processing are derived from Fourier analysis
- ▶ In this context, we assume that \mathbf{x} corresponds to the discrete measurement of a continuous signal $x(t)$
- ▶ **Sampling theory:** uniform sampling period T_s and sampling frequency

$$F_s = \frac{1}{T_s}$$

$$x[n] = x(nT_s)$$

Discrete Fourier Transform (DFT)

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-j2\pi \frac{kn}{N}} \text{ pour } 0 \leq k \leq N - 1$$

where N is the number of samples

- The space between two observable frequencies is called **frequency resolution**

$$\Delta f = \frac{F_s}{N}$$

- $X[k]$ corresponds to the DFT for the physical frequency


$$f[k] = k \frac{F_s}{N} \text{ for } 0 \leq k \leq N - 1$$

- No physical frequency greater than $\frac{F_s}{2}$ can be observed (Nyquist theorem).

Spectral analysis



- ▶ $X[k]$ is a complex quantity: most of the time, we use the squared absolute values $|X[k]|^2$ instead
- ▶ The analysis of the quantity $|X[k]|^2$ can allow to discover interesting properties of the time series
- ▶ $|X[k]|^2$ with low frequencies f_k correspond to phenomena with smooth variations
- ▶ $|X[k]|^2$ with large frequencies f_k correspond to phenomena with fast variations
- ▶ One very useful plot consists in plotting $|X[k]|^2$ as a function of $f[k]$: such plot is often referred to as **spectrum** (hence spectral analysis)

Example



Two classes of signals?

Example



Can be distinguished based on their DFT coefficients

Statistical vision

- ▶ In order to better understand the properties of a signal, deterministic analysis such as Fourier has been extended to probabilistic and statistical analysis
- ▶ In this context, we assume that $x[1 : N]$ corresponds to a realization of a stochastic process $X[1 : N]$
- ▶ Each $X[n]$ can be seen as a random variable
- ▶ Statistical properties of $X[1 : n]$ can be retrieved from estimates based on $x[1 : n]$

Two fundamental properties

- ▶ **Stationarity** : The statistical properties of the time series do not change over time

- ▶ Order 1

$$\forall n, \quad \mathbb{E}[X[n]] = \mu$$



- ▶ Order 2

$$\forall n_1, n_2, \quad \mathbb{E}[X[n_1]X[n_2]] = \gamma_X[|n_2 - n_1|]$$

- ▶ Order 1 + Order 2 → wide-sense stationarity (most common assumption)

- ▶ **Ergodicity** : Ensemble mean and the mean over time are equal (implies stationarity of order 1)

Stationarity vs. non-stationarity



- ▶ None of these signals are stationary
- ▶ In order to prevent this to happen, two solutions exist
 - ▶ Divide the signals into small frames where the signal is assumed to be stationary and ergodic (see later: spectrogram)
 - ▶ Use a change-point detection algorithm to detect these changes and work separately on each segment (see later: change-point detection)

Autocorrelation

Assuming that $X[1 : n]$ is ergodic and wide-sense stationary, we can estimate from $x[1 : n]$ the autocorrelation function


- ▶ Autocorrelation function

$$\hat{\gamma}_x^{biased}[m] = \frac{1}{N} \sum_{n=0}^{N-1} x[n]x[n+m]$$

where $x[n] = 0$ for $n \neq 0 \dots N - 1$


- ▶ This function helps (among other things) to discover the presence of periodic components within a signal

How to use the autocorrelation function




Original signal, sampling frequency 100 Hz

How to use the autocorrelation function



Autocorrelation function, peaks are visible for lags multiple of 50


How to use the autocorrelation function



A periodic signal with period $50 \times \frac{1}{100} = 0.5$ sec was hiding!

Spectrogram


- When the properties of the time series tend to change with time (non-stationary signals), it is more careful to compute the DFT on sliding windows
- By sliding the window along the signal, we recover a time-frequency representation called **spectrogram**



- Matrix representation: each column corresponds to the DFT on the window of interest.



X-axis: frame number, Y-axis: frequency bin

Example



Original signal, $F_s = 1000$ Hz



Example



Window length $N_w = 256$

Computation of the DFT on the first frame and storage in the spectrogram matrix...



Example



Window length $N_w = 256$



Computation of the DFT on the second frame and storage in the spectrogram matrix...

Example



Window length $N_w = 256$
Same process...

Example




Window length $N_w = 256$
Final result

DFT vs. Spectrogram



- ▶ Only use DFT when you are sure that there is no abrupt changes in the time series
- ▶ Note that using DFT will tend to average the frequency content on the whole time series, which can be tricky in some application contexts
- ▶ For safety, always first visualize the spectrogram to make sure that no significant changes occur

DFT vs. Spectrogram





Periodic phenomenon ? (sinusoid ?)

DFT vs. Spectrogram




DFT suggests a sinusoidal phenomenon around frequency 100 Hz

DFT vs. Spectrogram



In fact, chirp signal between 100 and 500 Hz !!

Hyperparameters for spectrogram




- ▶ N_w : window length (in samples)
Often taken as a power of 2 (for FFT) and linked to the desired frequency resolution.
- ▶ N_o : overlap between two successive frames (in samples)
Often taken as 50% or 75% of the window length and characterizes the time resolution (optimal when $N_o = N_w - 1$)
- ▶ w : analysis window (Hann, Hamming, Blackman...)
Traditionally, in order to limit side effects, the signal frame is multiplied by an analysis window

Contents

1. Introduction and data
2. Basic tools and notions for signal processing
3. Change-point detection
 - 3.1 Cost functions
 - 3.2 Search method
4. Convolutional dictionary learning

Problem 1: Change-Point Detection




Change-Point Detection

Given a time series x , retrieve the times (t_1, \dots, t_K) where a significant change occurs

- ▶ Necessitates to estimate both the change-points but also the number of changes K
- ▶ Highly depends on the meaning given to change


Problem statement



- ▶ When the changes are abrupt or when the estimation of the change-points is relevant in the context, we can use **change-point detection** methods
- ▶ Let assume that signal $x[n]$ undergoes abrupt changes at times
$$\mathcal{T}^* = (t_1^*, \dots, t_{K^*}^*)$$
- ▶ Goal: retrieve the number of change-points and K^* and their times \mathcal{T}^*
- ▶ One assumption: offline segmentation (but can easily be adapted to online setting) [Truong et al., 2020]

Problem statement

$$(\hat{t}_1, \dots, \hat{t}_K) = \operatorname{argmin}_{(t_1, \dots, t_K)} \sum_{k=0}^K c(x[t_k : t_{k+1}])$$




Cost function $c(.)$

- ▶ Measures the homogeneity of the segments
- ▶ Choosing $c(.)$ conditions the type of change-points that we want to detect
- ▶ Often based on a probabilistic model for the data

Problem solving

- ▶ Optimal resolution with dynamic programming
- ▶ Approximate resolution (sliding windows...)



Cost function

$$(\hat{t}_1, \dots, \hat{t}_K) = \operatorname{argmin}_{(t_1, \dots, t_K)} \sum_{k=0}^K c(x[t_k : t_{k+1}])$$

Convention : $t_0 = 0, t_{K+1} = N$

- ▶ Function $c(\cdot)$ is characteristic of the notion of *homogeneity*
- ▶ The most common cost functions are linked to parametric probabilistic models: in this case change-points are defined as changes in the parameters of a probability density function [Basseville et al., 1993]
- ▶ Non parametric cost functions can also be introduced when no model is available

Maximum likelihood estimation

Given a parametric family of distribution densities $f(\cdot|\theta)$ parametrized with $\theta \in \Theta$, a cost function can be derived:

$$c_{ML}(x[a : b]) = -\sup_{\theta} \sum_{n=a+1}^b \log f(x[n]|\theta)$$

- ▶ Corresponds to the assumption that on a regime, samples are i.i.d. according to a parametric distribution density
- ▶ On each regime, the parameters are estimated through maximum likelihood estimation
- ▶ This model can be adapted to several situations: change in mean, change in variance, change in both mean and variance...


Change in mean

The most popular is indubitably the L2 norm [Page, 1955]


$$c_{L_2}(x[a : b]) = \sum_{n=a+1}^b \|x[n] - \mu_{a:b}\|_2^2$$

where $\mu_{a:b}$ is the empirical mean of the segment $x[a : b]$.


- ▶ Particular case of c_{ML} with Gaussian model with fixed variance
- ▶ Allows to detect changes in mean




Example



Example: Change-Point Detection with c_{L_2}



Example: Change-Point Detection with c_{L_2}



Search method

$$(\hat{t}_1, \dots, \hat{t}_K) = \operatorname{argmin}_{(t_1, \dots, t_K)} \sum_{k=0}^K c(x[t_k : t_{k+1}])$$

Convention : $t_0 = 0$, $t_{K+1} = N$

- ▶ Several methods can be used to solve this problem with a fixed K
- ▶ Optimal resolution with dynamic programming: find the true solution of the problem (but costly)
- ▶ Approximated resolution with windows: test for one unique change-point on a window

Optimal resolution

- ▶ By denoting

$$V(\mathcal{T}, \mathbf{x}) = \sum_{k=0}^K c(x[t_k : t_{k+1}])$$

we can see that

$$\begin{aligned} \min_{|\mathcal{T}|=K} V(\mathcal{T}, \mathbf{x}) &= \min_{0=t_0 < t_1 < \dots < t_K < t_{K+1}=N} \sum_{k=0}^K c(x[t_k : t_{k+1}]) \\ &= \min_{t \leq T-K} \left[c(x[0 : t]) + \min_{t_0=t < t_1 < \dots < t_{K-1} < t_K=T} \sum_{k=0}^{K-1} c(x[t_k : t_{k+1}]) \right] \\ &= \min_{t \leq T-K} \left[c(x[0 : t]) + \min_{|\mathcal{T}|=K-1} V(\mathcal{T}, x[t : N]) \right] \end{aligned}$$

- ▶ Recursive problem: resolution with dynamic programming [Bai et al., 2003]
- ▶ Two steps: computation of the cumulative costs + determination of the change-points

Optimal resolution

Algorithm 1 Algorithm Opt

Input: signal $\{y_t\}_{t=1}^T$, cost function $c(\cdot)$, number of regimes $K \geq 2$.

for all (u, v) , $1 \leq u < v \leq T$ **do**

Initialize $C_1(u, v) \leftarrow c(\{y_t\}_{t=u}^v)$.

end for

for $k = 2, \dots, K - 1$ **do**

for all $u, v \in \{1, \dots, T\}$, $v - u \geq k$ **do**

$C_k(u, v) \leftarrow \min_{u+k-1 \leq t < v} C_{k-1}(u, t) + C_1(t+1, v)$

end for

end for

Initialize L , a list with K elements.

Initialize the last element: $L[K] \leftarrow T$.

Initialize $k \leftarrow K$.

while $k > 1$ **do**

$s \leftarrow L(k)$

$t^* \leftarrow \operatorname{argmin}_{k-1 \leq t < s} C_{k-1}(1, t) + C_1(t+1, s)$

$L(k-1) \leftarrow t^*$

$k \leftarrow k - 1$

end while

Remove T from L


Output: set L of estimated breakpoint indexes.

Complexity of $\mathcal{O}(KN^2)$



Contents

1. Introduction and data
2. Basic tools and notions for signal processing
3. Change-point detection
4. Convolutional dictionary learning
 - 4.1 Problem
 - 4.2 Alternated resolution

Problem 2: Pattern Extraction



Input time series



Extracted patterns

Problem 2: Pattern Extraction

Pattern Extraction

Given an input time series x (or a set of time series), learn a dictionary of patterns \mathcal{P}

- ▶ A template is a *shape* that appear repetitively in the time series (but kinda blurry notion)
- ▶ All templates are supposed to have the same length (for sake of simplicity)
- ▶ The extracted patterns can be used to characterize the time series, or studied individually



Dictionary-based pattern extraction

- ▶ Finding patterns in a time series can be seen as a dictionary learning optimization problem
- ▶ Given an input time series \mathbf{x} , learn a dictionary of K patterns \mathbf{d}_k of length L
- ▶ These patterns can be activated : activations \mathbf{z}_k of length $N - L + 1$

$z_k[n] \neq 0$ if pattern \mathbf{d}_k is activated at time n

[Grosse et al., 2007 ; Wohlberg, 2014]

Convolutional dictionary learning



Convolutional dictionary learning

Given a time series \mathbf{x} , number of pattern K and pattern length L , learn

- ▶ Patterns \mathbf{d}_k of length L
- ▶ Activation signals \mathbf{z}_k of length $N - L + 1$

$$\mathbf{x}[n] = \sum_{k=1}^K (\mathbf{z}_k * \mathbf{d}_k)[n] + e[n]$$

Optimization problem

$$\min_{(\mathbf{d}_k), (\mathbf{z}_k)} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2 + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

$\forall k, \|\mathbf{d}_k\|_2^2 \leq 1$

- ▶ **Normalization constraint** for the dictionary atoms \mathbf{d}_k , that prevents numerical instabilities (otherwise setting $\alpha \mathbf{d}_k$ and $\alpha^{-1} \mathbf{z}_k$ gives the same result)
- ▶ **Sparsity constraint** for the activations \mathbf{z}_k , that improves the interpretability of the learned patterns (see Lecture 3)

Not convex with respect to the couple $(\mathbf{d}_k), (\mathbf{z}_k)$ but convex when the subproblems are taken individually

Alternated resolution

Alternated resolution of two subproblems

Dictionary learning

$$\mathbf{D}^* = \underset{\substack{\mathbf{D} = (\mathbf{d}_1, \dots, \mathbf{d}_K) \\ \forall k, \|\mathbf{d}_k\|_2^2 \leq 1}}{\operatorname{argmin}} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2$$

Convolutional sparse coding

$$\mathbf{Z}^* = \underset{\mathbf{Z} = (\mathbf{z}_1, \dots, \mathbf{z}_K)}{\operatorname{argmin}} \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2 + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

Alternated resolution

$$\min_{\substack{(\mathbf{d}_k), (\mathbf{z}_k) \\ \forall k, \|\mathbf{d}_k\|_2^2 \leq 1}} \underbrace{\left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2}_{f(\mathbf{Z}, \mathbf{D})} + \lambda \sum_{k=1}^K \|\mathbf{z}_k\|_1$$

- ▶ Both these problems can be solved with Proximal Gradient Descent algorithms
- ▶ Two main steps :
 1. Gradient descent step w.r.t. $\nabla_{\mathbf{D}} f(\mathbf{Z}, \mathbf{D})$ or $\nabla_{\mathbf{Z}} f(\mathbf{Z}, \mathbf{D})$
 2. Proximal step to *project* the update on the constraint set
- ▶ The main question is therefore: how can we easily compute the gradients?

Formulation of the gradients

$$f(\mathbf{Z}, \mathbf{D}) = \left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2$$

- ▶ By using the convolution theorem and the Parseval theorem, we have that

$$f(\mathbf{Z}, \mathbf{D}) = \left\| \hat{\mathbf{x}} - \sum_{k=1}^K \widehat{\mathbf{z}}_k \odot \widehat{\mathbf{d}}_k \right\|_2^2$$

where $\widehat{\cdot}$ denotes the Discrete Fourier Transform (DFT) and \odot the component-wise product

- ▶ As $\mathbf{u} \odot \mathbf{d} = \text{diag}(\mathbf{u})\mathbf{v}$, this expression can be rewritten in the matrix form as

$$f(\mathbf{Z}, \mathbf{D}) = \left\| \hat{\mathbf{x}} - \widehat{\mathbf{D}} \hat{\mathbf{z}} \right\|_2^2$$

- ▶ The gradient of this quantity is easy to compute w.r.t. $\widehat{\mathbf{D}}$ and $\hat{\mathbf{z}}$
- ▶ Updates for \mathbf{D} and \mathbf{Z} can then be estimated by performing inverse DFT

Proximal operators

- ▶ For the unit ball constraint (dictionary), the proximal operator writes as

$$\text{proj}_{\|\cdot\|_2^2 \leq 1}(\mathbf{y}) = \frac{\mathbf{y}}{\max(\|\mathbf{y}\|_2^2, 1)}$$

Projection on the unit ball

- ▶ For the L1-sparsity constraint (atoms), the proximal operator writes as

$$\mathcal{S}\gamma(\mathbf{y})[n] = \text{sign}(y[n]) \times \max(|y[n]| - \gamma, 0)$$

Soft thresholding operator

Dictionary learning

$$\mathbf{D}^* = \underset{\substack{\mathbf{D}=(\mathbf{d}_1, \dots, \mathbf{d}_K) \\ \forall k, \|\mathbf{d}_k\|_2^2 \leq 1}}{\operatorname{argmin}} \left\| \mathbf{x} - \underbrace{\sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k}_{f(\mathbf{Z}, \mathbf{D})} \right\|_2^2$$

- ▶ Basic approach : **Proximal Gradient Descent** with fixed \mathbf{Z}

1. Gradient step :

$$\mathbf{D} \leftarrow \mathbf{D} - \alpha \nabla_{\mathbf{D}} f(\mathbf{Z}, \mathbf{D})$$

2. Proximal projection step :

$$\mathbf{d}_k \leftarrow \text{proj}_{\|\cdot\|_2^2 \leq 1}(\mathbf{d}_k) = \frac{\mathbf{d}_k}{\max(\|\mathbf{d}_k\|_2^2, 1)}$$


- ▶ Other approaches : Alternate Direction Method of Multiplier (ADMM), K-SVD... (see [Mairal et al., 2010] and Lecture 3 for more details)

Convolutional sparse coding


$$\mathbf{Z}^* = \underset{\mathbf{Z}=(\mathbf{z}_1, \dots, \mathbf{z}_K)}{\operatorname{argmin}} \underbrace{\left\| \mathbf{x} - \sum_{k=1}^K \mathbf{z}_k * \mathbf{d}_k \right\|_2^2}_{f(\mathbf{Z}, \mathbf{D})} + \lambda \underbrace{\sum_{k=1}^K \|\mathbf{z}_k\|_1}_{\psi(\mathbf{Z})}$$

- ▶ Basic approach : **Iterative Soft Thresholding Algorithm (ISTA)** with fixed \mathbf{D} [Daubechies et al., 2004]
 1. Gradient step
$$\mathbf{Z} \leftarrow \mathbf{Z} - \alpha \nabla_{\mathbf{Z}} f(\mathbf{Z}, \mathbf{D})$$
 2. Proximal step
$$\mathbf{Z} = \mathcal{S}_{\lambda\alpha}(\mathbf{Z})$$
- ▶ Other approaches: Alternate Direction Method of Multiplier (ADMM), Fast Iterative Soft Thresholding Algorithm (FISTA), Coordinate Descent (CD) (see [Mairal et al., 2010] and Lecture 3 for more details)


Results




Results




Results




Results




Results




Results



Results




Results




One pattern

Results



Two patterns

Results



Mixture of patterns

Summary

- ▶ CDL allows to detect mixture of patterns thanks to the additive model
- ▶ CDL is traditionally quite sensitive to initialization and/or hyperparameters : random initializations (normalized gaussian noise), randomly chosen data frames, Lipschitz constant for gradient descent...
- ▶ CDL can also be used to search patterns from a fixed dictionary of templates : in this case, only perform the convolutional sparse coding step